

Suricata

Website monitor

Developed by Simon Brandeis

Starting the software

Prerequisites

Go must be installed to build the project golang.org.

Execution instructions

`go build` to build the executable, then just execute the file "suricata"

You can also execute the bash script `build_start` to build and launch "suricata".

Configuration

The websites to monitor and the check intervals are defined in a config file.

By default, "suricata" uses `config.sample`, but you can specify your own file by passing its path to the flag `cfg`.

Ex: `./suricata -cfg="/your/config/file"`

Config file format:

```
www.google.com,300
https://github.com,500
```

This config file sets suricata to monitor google.com and github.com with a check interval of respectively 300 and 500 milliseconds

Documentation

Folder structure

```
| -main.go
|  config.sample
| -suricata
| -README.md
| -bin
|   | -build_start
| -cui
|   | -Ui.go
```

```
| -monitor
|   | -Aggregator_test.go
|   | -MaxHeap.go
|   | -Aggregator.go
|   | -MaxHeap_test.go
|   | -Pinger.go
|   | -Orchestrator.go
|   | -Orchestrator_test.go
|   | -Alert.go
|   | -Report.go
|   | -Website.go
```

"suricata" is composed of two modules: - suricata/monitor which monitors the websites - suricata/cui which abstracts UI updating.

"suricata" has 1 external dependency: termui.

Functional Overview

Website monitoring is centralized by the singleton object Orchestrator.

Orchestrator registers websites (i.e, urls and check intervals) and controls the Pinger objects, which perform the HTTP requests.

Pinger objects perform requests in parallel, and send data about the response in a shared channel (pipeline in main.go). This data is represented by a PingLog object.

When data is sent to the pipeline, Orchestrator forwards it to the appropriate Aggregator objects.

Aggregator objects perform on-the-fly metrics computing. To do so, PingLog are stored in a queue, and Aggregator holds a reference to the first and last elements of this queue.

Aggregators.png illustrates the data structure.



Metrics computing is done when aggregating (queueing) or de-aggregating (dequeueing) a PingLog. Queueing and dequeueing have a $O(1)$ complexity.

Plus, Aggregator holds a reference to each element of the queue in a MaxHeap attribute, backed by an array. Each element of the queue holds its index in this array. Elements are ordered by their ResponseTime attribute.

This allows to insert and delete elements in the heap in $O(\ln(n))$, where n is the number of elements in the heap. Most importantly, this data structure allows maximum element retrieval in $O(1)$ time complexity.

Eventually, processing the incoming PingLogs is done in $O(\ln(n))$ time complexity, where n is the number of elements in the heap, and yields average and maximum values on the data processed.

Metrics on aggregated logs are regularly read to update Report objects, in which metrics are stored. Reports are transformed into Summary, used by the the suricata/cui package to generate a termui.Table components to display on screen.

During the process, Alerts objects are emitted on the alert channel, transporting either messages or serious alerts (eg, when a website's availability drops under 80%).

Possible Improvements

For the same website, all associated Aggregators objects hold the reference to the latest PingLog. It would be better to have them share this reference: it's less bug-prone and slightly more efficient.

The combination of a Max Heap and a Linked List seems a bit clumsy, it feels like we store twice too much data.

If computing the maximum response time is not needed, dropping the heap improves significantly the performance: insertion and deletion of PingLogs is in $O(1)$ time complexity.

This tool is designed to be running indefinitely, it should have a logging system to store alerts and measures somewhere on the file system: we can't just spend the day staring at it.

Test coverage is only of 49 %. TDD could have been useful to highlight bugs earlier.

Orchestrator.go is a quite long file (~340 lines), which affects readability . Perhaps it could be splitted in smaller files.